
Innovations in Modern Database Systems

Predictive Modeling, Generative AI, and Serverless Co-Design

CSCE – 5370 Distributive & Parallel Database Systems

Gayathri Baman

ID: 11697946

OUTLINE

- Introduction
- Descriptive Terminology for the Three Models
- Strengths and Weakness for the Models
- Comparative Analysis
- References

Introduction

- The modern database workload environment is, therefore, not only extremely heterogeneous but also distributed and dynamic. The requirements from the applications include accurate performance predictions, low latency, highly-scalable multi-tenancy, and easy usability for even a non-expert user. All these have led to an evolving change in the architecture of databases from intelligent and adaptive to that of modular design.

Why Are New Database Concepts Needed?



Old database systems were for simpler sets of requirements



New requirements:



Instant Answers-(Low latency)



Everyone Must Use It- User Friendliness



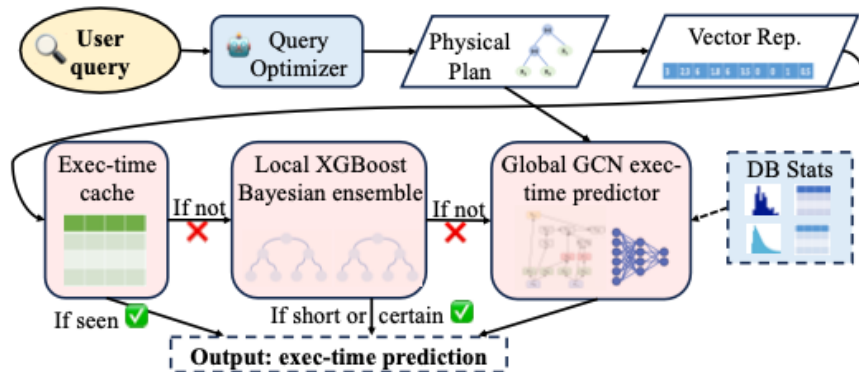
Many Users Simultaneously- Scalability



New tools are required to hear these demands.

Stage: Stage — Predicting Query Time in Amazon Redshift

Stage: Query Execution Time Prediction in Amazon Redshift



- Purpose: Predict SQL query execution time in Amazon Redshift
- Multi-layer system:
 - Cache : If seen before → use previous time
 - Local XGBoost Ensemble : Fast machine learning guess
 - Global Graph Neural Network (GNN) : Deep learning for new queries

Pro's & Con's

Strengths:

Good performance in short and medium queries,

Fast inference for repeated queries (cache/local).

Generally, falls back with GNN while failing generalization provisions.

Weaknesses:

GNN inference latency (~100ms).

Failed for new unseen long queries with no generalization.

Requires lots of memory resources for operations and training.

Verification of Relational Database Languages Codes Generated by ChatGPT

- Can write SQL, relational algebra (RA), and relational calculus (RC)
- Learns from natural language input
- Good for learning and quick coding

Example: “Find vendor names in Bangkok”

- SQL: Simple command -> `SELECT VNAME FROM VENDOR WHERE LOCATION = 'Bangkok';`
- RA: Uses selection and projection -> $\Pi_VNAME (\sigma_location = 'Bangkok'(VENDOR))$
- RC: Uses logical expression -> $v. VNAME \mid v \in VENDOR \wedge v. Location = 'Bangkok'.$

In this example, relational calculus queries have reflected formal logic through them existential and universal quantifiers, demonstrating a great capability beyond the casual syntax of SQL.

Pro's & Con's

Strengths:

All relational operators are supported.

Formulates correct queries in SQL, RA, and RC.

Very lightweight, User friendly.

Weaknesses:

Semantic superficiality or schema-specific optimization may also be missing.

Can get confused with tricky schemas

Doesn't optimize queries or aware of indexes.

Serverless Runtime / Database Co-Design With Asynchronous I/O

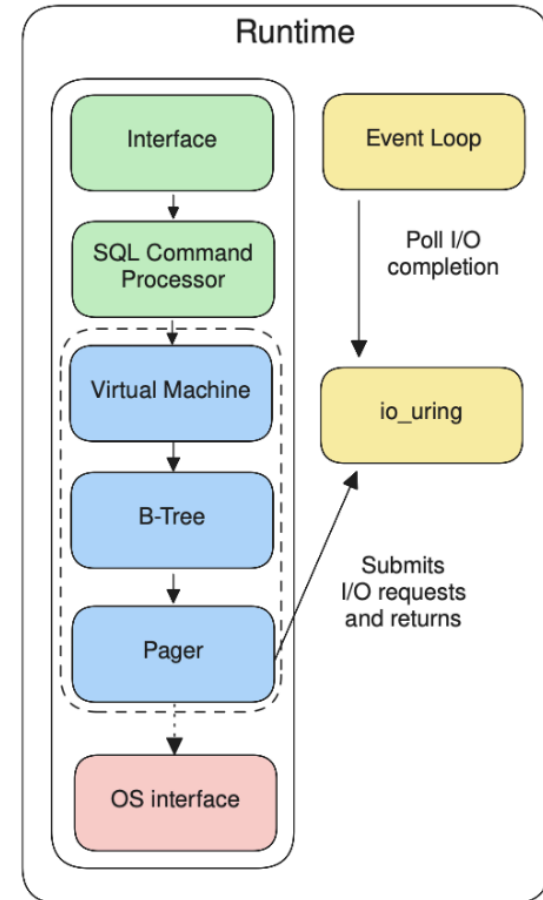
Limbo modifies SQLite for serverless runtimes by replacing blocking bytecode instructions (e.g., Next) with asynchronous versions.

How it works:

- Next → NextAsync → NextWait
- Separates query logic from disk access
- system handle many queries at once.

Results:

- Simulated 100 users querying `SELECT * FROM users`
- Limbo: stays fast and consistent



Pro's & Con's

Strengths:

Up to 100x tail latency reduced.

Cohorts can handle massive concurrent workloads.

Very efficient CPU and I/O resource usage.

Weaknesses:

With runtime integration it's a real hassle.

Clean-slate design isn't compatible with SQLite.

Requires runtime support for `io_uring` and event loops.

Comparative Analysis

Criteria	Stage	ChatGPT Code Verification	Limbo
Speed	✅ Fast (cache/local); ❌ Slower GNN (~100ms)	✅ Instant generation	✅ Runtime latency reduction up to 100x
Accuracy	✅ High for short/medium; ❌ Uncertain on long queries	✅ Correct across SQL/RA/RC	✅ Accurate with async logic
Memory Usage	❌ Higher (ensemble + GNN)	✅ Minimal (prompt-based)	✅ Efficient (coroutines + I/O buffer)
Scalability	⚠️ Limited by model cost	✅ Scales with demand	✅ Excellent (supports many tenants)
Robustness	✅ Uses uncertainty-aware switching	⚠️ Limited schema understanding	✅ Decoupled, modular architecture
Usability	⚠️ Backend-only; not user-facing	✅ Ideal for learners and non-programmers	⚠️ Developer-level usage only
Optimization Awareness	✅ Avoids unnecessary model calls	❌ No query performance optimization	✅ Asynchronous model avoids thread waste
Implementation Effort	❌ High (training, integration)	✅ None for users; ⚠️ Hard to verify semantics manually	❌ High (requires system-level engineering)

References

- [1] D. J. Abadi, D. Carney, U. Çetintemel, M. Cherniack, C. Convey, S. Lee, M. Stonebraker, N. Tatbul, and S. Zdonik. 2003. Aurora: a data stream management system. In *Proceedings of the 2003 ACM SIGMOD international conference on Management of data* (SIGMOD '03). Association for Computing Machinery, New York, NY USA, 666. <https://doi.org/10.1145/872757.872858>
- [2] Ziniu Wu, Alex Ratner, Cong Yu, and Matei Zaharia. 2024. Stage: Query Execution Time Prediction in Amazon Redshift. In *Proceedings of the 2024 International Conference on Management of Data (SIGMOD-Companion '24)*. Association for Computing Machinery, New York, NY , USA. <https://doi.org/10.1145/3626246.3653391>
- [3] Putsadee Pornphol and Suphamit Chittayasothorn. 2023. Verification of Relational Database Languages Codes Generated by ChatGPT. In *Proceedings of the 2023 Asia Symposium on Software Engineering (ASSE 2023)*. Association for Computing Machinery, New York, NY , USA. <https://doi.org/10.1145/3634814.3634817>

THANK YOU
